

Project Title : **Flight Delay Prediction**

Aim :

The aim of the project is to predict if the flight is Delayed or not

Team Members:

Pragadeeshwar N E0119003

Aswin G E0119062

Lohit Arun S E0119042

▼ Importing libraries

```
import time
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import warnings
warnings.filterwarnings("ignore")

from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, VotingClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split

from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import plot_confusion_matrix
from sklearn.metrics import precision_recall_fscore_support as score
from sklearn.metrics import auc

import tensorflow as tf

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive
```

▼ Importing the dataset

```
flights = pd.read_csv('/content/drive/MyDrive/flights.csv')
flights
```

	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT
0	2015	1	1	4	AS	98	N407AS	
1	2015	1	1	4	AA	2336	N3KUAA	
2	2015	1	1	4	US	840	N171US	
3	2015	1	1	4	AA	258	N3HYAA	
4	2015	1	1	4	AS	135	N527AS	
...
5819074	2015	12	31	4	B6	688	N657JB	
5819075	2015	12	31	4	B6	745	N828JB	
5819076	2015	12	31	4	B6	1503	N913JB	
5819077	2015	12	31	4	B6	333	N527JB	
5819078	2015	12	31	4	B6	839	N534JB	

5819079 rows × 31 columns



▼ Data Preprocessing

```
flights_needed_data = flights[0:100000]
flights_needed_data
```


	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT
0	2015	1	1	4	AS	98	N407AS	
1	2015	1	1	4	AA	2336	N3KUAA	
2	2015	1	1	4	US	840	N171US	
3	2015	1	1	4	AA	258	N3HYAA	
4	2015	1	1	4	AS	105	N507AS	

```
flights_needed_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   YEAR                                100000 non-null  int64
1   MONTH                              100000 non-null  int64
2   DAY                                100000 non-null  int64
3   DAY_OF_WEEK                        100000 non-null  int64
4   AIRLINE                            100000 non-null  object
5   FLIGHT_NUMBER                      100000 non-null  int64
6   TAIL_NUMBER                        99833 non-null   object
7   ORIGIN_AIRPORT                    100000 non-null  object
8   DESTINATION_AIRPORT              100000 non-null  object
9   SCHEDULED_DEPARTURE              100000 non-null  int64
10  DEPARTURE_TIME                    97702 non-null   float64
11  DEPARTURE_DELAY                   97702 non-null   float64
12  TAXI_OUT                          97629 non-null   float64
13  WHEELS_OFF                       97629 non-null   float64
14  SCHEDULED_TIME                    100000 non-null  float64
15  ELAPSED_TIME                      97387 non-null   float64
16  AIR_TIME                          97387 non-null   float64
17  DISTANCE                          100000 non-null  int64
18  WHEELS_ON                        97560 non-null   float64
19  TAXI_IN                          97560 non-null   float64
20  SCHEDULED_ARRIVAL                100000 non-null  int64
21  ARRIVAL_TIME                     97560 non-null   float64
22  ARRIVAL_DELAY                    97387 non-null   float64
23  DIVERTED                         100000 non-null  int64
24  CANCELLED                        100000 non-null  int64
25  CANCELLATION_REASON              2389 non-null    object
26  AIR_SYSTEM_DELAY                 34625 non-null   float64
27  SECURITY_DELAY                   34625 non-null   float64
28  AIRLINE_DELAY                    34625 non-null   float64
29  LATE_AIRCRAFT_DELAY              34625 non-null   float64
30  WEATHER_DELAY                    34625 non-null   float64
dtypes: float64(16), int64(10), object(5)
memory usage: 23.7+ MB
```

```
sns.jointplot(data=flights_needed_data, x="SCHEDULED_ARRIVAL", y="ARRIVAL_TIME")
```

```
<seaborn.axisgrid.JointGrid at 0x7f8d617cf390>
```

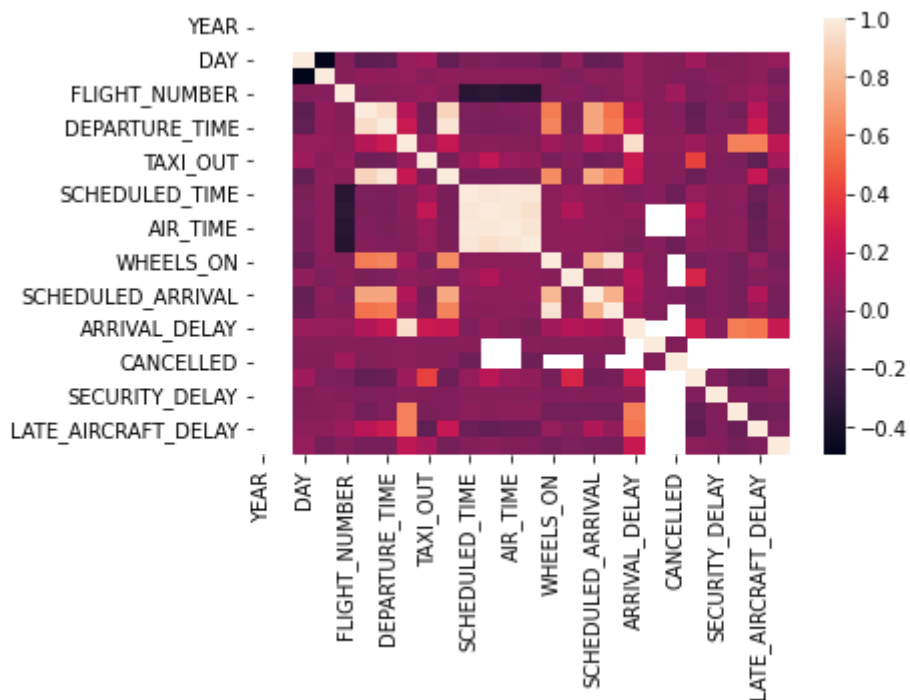


```
corr = flights_needed_data.corr(method='pearson')
```

SCHEDULED ARRIVAL

```
sns.heatmap(corr)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8d57d6da50>
```



```
corr
```

	YEAR	MONTH	DAY	DAY_OF_WEEK	FLIGHT_NUMBER	SCHEDULED
YEAR	NaN	NaN	NaN	NaN	NaN	
MONTH	NaN	NaN	NaN	NaN	NaN	
DAY	NaN	NaN	1.000000	-0.497084	0.004412	
DAY_OF_WEEK	NaN	NaN	-0.497084	1.000000	0.010955	
FLIGHT_NUMBER	NaN	NaN	0.004412	0.010955	1.000000	
SCHEDULED_DEPARTURE	NaN	NaN	-0.138130	0.046914	-0.003027	
DEPARTURE_TIME	NaN	NaN	-0.124369	0.045182	0.010140	
DEPARTURE_DELAY	NaN	NaN	0.060064	0.055632	0.034863	
TAXI_OUT	NaN	NaN	0.093451	0.007291	0.061010	
WHEELS_OFF	NaN	NaN	-0.119781	0.044150	0.016377	
SCHEDULED_TIME	NaN	NaN	-0.026285	0.019755	-0.337801	
ELAPSED_TIME	NaN	NaN	-0.018470	0.029025	-0.318819	
AIR_TIME	NaN	NaN	-0.036330	0.030678	-0.339135	
DISTANCE	NaN	NaN	-0.035208	0.024666	-0.356196	
WHEELS_ON	NaN	NaN	-0.095731	0.013749	-0.003670	
TAXI_IN	NaN	NaN	0.037407	-0.017789	0.014464	
SCHEDULED_ARRIVAL	NaN	NaN	-0.110820	0.031725	-0.018891	
ARRIVAL_TIME	NaN	NaN	-0.091687	0.011477	0.000753	
ARRIVAL_DELAY	NaN	NaN	0.070770	0.067520	0.056163	
DIVERTED	NaN	NaN	0.004847	-0.000709	0.007155	
CANCELLED	NaN	NaN	-0.006000	-0.006409	0.090008	
AIR_SYSTEM_DELAY	NaN	NaN	0.097693	-0.019626	-0.032564	
SECURITY_DELAY	NaN	NaN	-0.010550	0.008156	-0.007260	
AIRLINE_DELAY	NaN	NaN	-0.001603	0.003648	0.023770	
LATE_AIRCRAFT_DELAY	NaN	NaN	0.033213	0.033729	0.076581	

```
# filtering out unnecessary columns
```

```
flights_needed_data=flights_needed_data.drop(['YEAR','FLIGHT_NUMBER','AIRLINE','DISTANCE',
                                                'SCHEDULED_TIME','DEPARTURE_TIME','WHEELS_OF',
                                                'AIR_TIME','WHEELS_ON','DAY_OF_WEEK','TAXI_I',
                                                axis=1])
```

```
flights_needed_data
```

	MONTH	DAY	ORIGIN_AIRPORT	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE	DEPARTU
0	1	1	ANC	SEA	5	
1	1	1	LAX	PBI	10	
2	1	1	SFO	CLT	20	
3	1	1	LAX	MIA	20	
4	1	1	SEA	ANC	25	
...
99995	1	7	ATL	BQK	1108	
99996	1	7	LAS	PHL	1108	
99997	1	7	SFO	BFL	1108	
99998	1	7	ORD	MCO	1109	
99999	1	7	HOU	DFW	1109	

100000 rows × 16 columns

```
# replacing all NaN values with the mean of the attribute in which they are present
flights_needed_data=flights_needed_data.fillna(flights_needed_data.mean())
flights_needed_data
```

	MONTH	DAY	ORIGIN_AIRPORT	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE	DEPARTU
0	1	1	ANC	SEA	5	
1	1	1	LAX	PBI	10	
2	1	1	SFO	CLT	20	
3	1	1	LAX	MIA	20	
4	1	1	SEA	ANC	25	
...
99995	1	7	ATL	BQK	1108	
99996	1	7	LAS	PHL	1108	
99997	1	7	SFO	BFL	1108	
99998	1	7	ORD	MCO	1109	
99999	1	7	HOU	DFW	1109	

100000 rows × 16 columns



```
# creating a new column; it will tell if the flight was delayed or not
```

```

result=[]

for row in flights_needed_data['ARRIVAL_DELAY']:
    if row > 15:
        result.append(1)
    else:
        result.append(0)

```

```
flights_needed_data['result'] = result
```

```
flights_needed_data
```

	MONTH	DAY	ORIGIN_AIRPORT	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE	DEPARTU
0	1	1	ANC	SEA	5	
1	1	1	LAX	PBI	10	
2	1	1	SFO	CLT	20	
3	1	1	LAX	MIA	20	
4	1	1	SEA	ANC	25	
...	
99995	1	7	ATL	BQK	1108	
99996	1	7	LAS	PHL	1108	
99997	1	7	SFO	BFL	1108	
99998	1	7	ORD	MCO	1109	
99999	1	7	HOU	DFW	1109	

100000 rows × 17 columns



```
flights_needed_data.value_counts('result')
```

```

result
0    63779
1    36221
dtype: int64

```

```
# removing some more columns
```

```

flights_needed_data=flights_needed_data.drop(['ORIGIN_AIRPORT', 'DESTINATION_AIRPORT', 'AR
flights_needed_data

```

	MONTH	DAY	SCHEDULED_DEPARTURE	DEPARTURE_DELAY	SCHEDULED_ARRIVAL	DIVERTED
0	1	1	5	-11.0	430	0
1	1	1	10	-8.0	750	0
2	1	1	20	-2.0	806	0
3	1	1	20	-5.0	805	0
4	1	1	25	-1.0	320	0
...
99995	1	7	1108	-6.0	1219	0
99996	1	7	1108	9.0	1842	0
99997	1	7	1108	-7.0	1225	0
99998	1	7	1109	7.0	1454	0
99999	1	7	1109	-9.0	1220	0

▼ Splitting the dataset into Train and Test data

```
data = flights_needed_data.values
X, y = data[:, :-1], data[:, -1]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=202)

# define a function to train and test the models
def train_model(model, X_train, y_train):
    model.fit(X_train, y_train)
    return model

def test_model(model, X_test, y_test):
    predictions = model.predict(X_test)
    accuracy = model.score(X_test, y_test)
    metrics_report = classification_report(y_test, predictions)
    precision, recall, fscore, train_support = score(y_test, predictions, average='weighted')
    return predictions, accuracy, metrics_report, (precision, recall, fscore)
```

▼ Feature Scaling : Standardisation

```
scaled_features = StandardScaler()
X_train = scaled_features.fit_transform(X_train)
X_test = scaled_features.transform(X_test)
```

▼ C01 : Base Learners

Decision Tree :

```
# define the model
dt_model = DecisionTreeClassifier(random_state=42)

# fit the model
dt_model = train_model(dt_model, X_train, y_train)
print(dt_model)

DecisionTreeClassifier(random_state=42)
```

Logistic Regression:

```
# define the model
log_model = LogisticRegression(penalty='l2', max_iter=500)

# fit the model
log_model= train_model(log_model, X_train,y_train)
print(log_model)

LogisticRegression(max_iter=500)
```

Support Vector Machine:

```
svm_model= SVC(kernel = 'linear', random_state = 0)
svm_model = train_model(svm_model,X_train,y_train)
print(svm_model)

SVC(kernel='linear', random_state=0)
```

▼ C02 :Ensemble Learning :**Hard Voting**

```
estimators = [('Decision Tree',dt_model),('Logistic Regression',log_model),('Support Vecto

ensemble_model = VotingClassifier(estimators=estimators,voting='hard')
ensemble_model = train_model(ensemble_model,X_train,y_train)
```

Gradient Boosting

```
# define the model
gb_model = GradientBoostingClassifier(n_estimators=50, max_depth=10, random_state=2)
```

```
# fit the model
gb_model = train_model(gb_model, X_train,y_train)
print(gb_model)

GradientBoostingClassifier(max_depth=10, n_estimators=50, random_state=2)
```

▼ C03 : Dimensionality Reduction

▼ *Principal Component Analysis*

```
from sklearn.decomposition import PCA
pca = PCA(n_components=9)
X1_train = pca.fit_transform(X_train)
X1_test = pca.transform(X_test)

# define the model
log_model_pca = LogisticRegression(penalty='l2', max_iter=500)

# fit the model
log_model_pca= train_model(log_model_pca, X1_train,y_train)
print(log_model_pca)

LogisticRegression(max_iter=500)
```

▼ C04 : Artificial Neural Network

```
ann = tf.keras.models.Sequential()

ann.add(tf.keras.layers.Dense(units=10, activation='relu'))

ann.add(tf.keras.layers.Dense(units=10, activation='relu'))

ann.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))

ann.compile(optimizer = 'adam',loss = 'binary_crossentropy' ,metrics = ['accuracy'])

ann.fit(X1_train, y_train, batch_size = 32, epochs = 100)

Epoch 75/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0039 - accuracy
Epoch 74/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0032 - accuracy
Epoch 75/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0032 - accuracy
Epoch 76/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0032 - accuracy
```

```
2500/2500 [=====] - 5s 2ms/step - loss: 0.0041 - accuracy
Epoch 77/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0027 - accuracy
Epoch 78/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0037 - accuracy
Epoch 79/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0028 - accuracy
Epoch 80/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0034 - accuracy
Epoch 81/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0031 - accuracy
Epoch 82/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0037 - accuracy
Epoch 83/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0028 - accuracy
Epoch 84/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0029 - accuracy
Epoch 85/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0039 - accuracy
Epoch 86/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0021 - accuracy
Epoch 87/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0035 - accuracy
Epoch 88/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0025 - accuracy
Epoch 89/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0036 - accuracy
Epoch 90/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0032 - accuracy
Epoch 91/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0041 - accuracy
Epoch 92/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0040 - accuracy
Epoch 93/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0039 - accuracy
Epoch 94/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0032 - accuracy
Epoch 95/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0033 - accuracy
Epoch 96/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0026 - accuracy
Epoch 97/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0041 - accuracy
Epoch 98/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0020 - accuracy
Epoch 99/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0043 - accuracy
Epoch 100/100
2500/2500 [=====] - 5s 2ms/step - loss: 0.0020 - accuracy
<keras.callbacks.History at 0x7f8d40569650>
```

▼ C05: Comparison of Models

Decision Tree Prediction:

```
# Predicting the Test set results
from sklearn.metrics import accuracy_score
y_pred = dt_model.predict(X_test)
dt_accuracy = accuracy_score(y_pred,y_test)
metrics_report = classification_report(y_test, y_pred)
precision, recall, fscore, train_support = score(y_test, y_pred, average='weighted')
dt_prf =(precision, recall, fscore)
print('accuracy: {}'.format(dt_accuracy))
print('='*100)
print(metrics_report)
plot_confusion_matrix(dt_model,X_test,y_test,xticks_rotation='vertical', cmap="OrRd")
```


accuracy: 0.9981

```
=====
              precision    recall  f1-score   support

   0.0         1.00      1.00      1.00     12704
   1.0         1.00      1.00      1.00      7296

 accuracy          1.00      1.00      1.00     20000
 macro avg         1.00      1.00      1.00     20000
weighted avg         1.00      1.00      1.00     20000
```

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f8d4046c950>



Support Vector Machine Prediction:


```
# Predicting the Test set results
from sklearn.metrics import accuracy_score
y_pred = svm_model.predict(X_test)
svm_accuracy = accuracy_score(y_pred,y_test)
metrics_report = classification_report(y_test, y_pred)
precision, recall, fscore, train_support = score(y_test, y_pred, average='weighted')
svm_prf =(precision, recall, fscore)
print('accuracy: {}'.format(svm_accuracy))
print('='*100)
print(metrics_report)
plot_confusion_matrix(svm_model,X_test,y_test,xticks_rotation='vertical', cmap="Blues")
```

accuracy: 0.98605

```
=====
```

	precision	recall	f1-score	support
0.0	1.00	0.98	0.99	12704
1.0	0.97	0.99	0.98	7296
accuracy			0.99	20000
macro avg	0.98	0.99	0.99	20000
weighted avg	0.99	0.99	0.99	20000

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f8d55f0e5d0>



Logistic Regression Prediction

```
# Predicting the Test set results
from sklearn.metrics import accuracy_score
y_pred = log_model.predict(X_test)
log_accuracy = accuracy_score(y_pred,y_test)
metrics_report = classification_report(y_test, y_pred)
precision, recall, fscore, train_support = score(y_test, y_pred, average='weighted')
log_prf =(precision, recall, fscore)
print('accuracy: {}'.format(log_accuracy))
print('='*100)
print(metrics_report)
plot_confusion_matrix(log_model,X_test,y_test,xticks_rotation='vertical', cmap="GnBu")
```


accuracy: 0.9848

	precision	recall	f1-score	support
0.0	1.00	0.98	0.99	12704
1.0	0.97	0.99	0.98	7296
accuracy			0.98	20000
macro avg	0.98	0.99	0.98	20000
weighted avg	0.99	0.98	0.98	20000

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f8d402c3310>



Logistic Regression after using PCA




```
# Predicting the Test set results
from sklearn.metrics import accuracy_score
y_pred = log_model_pca.predict(X1_test)
log_pca_accuracy = accuracy_score(y_pred,y_test)
metrics_report = classification_report(y_test, y_pred)
precision, recall, fscore, train_support = score(y_test, y_pred, average='weighted')
log_pca_prf =(precision, recall, fscore)
print('accuracy: {}'.format(log_pca_accuracy))
print('='*100)
print(metrics_report)
plot_confusion_matrix(log_model_pca,X1_test,y_test,xticks_rotation='vertical', cmap="GnBu")
```

accuracy: 0.91315

	precision	recall	f1-score	support
0.0	0.89	0.99	0.94	12704
1.0	0.97	0.79	0.87	7296
accuracy			0.91	20000
macro avg	0.93	0.89	0.90	20000
weighted avg	0.92	0.91	0.91	20000

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f8d404d0dd0>



ANN Prediction

```
# Predicting the Test set results
from sklearn.metrics import accuracy_score
y_pred = ann.predict(X1_test)
y_pred = (y_pred > 0.5)
ann_accuracy = accuracy_score(y_pred,y_test)
metrics_report = classification_report(y_test, y_pred)
precision, recall, fscore, train_support = score(y_test, y_pred, average='weighted')
ann_prf =(precision, recall, fscore)
print('accuracy: {}'.format(ann_accuracy))
print('='*100)
print(metrics_report)
cm= confusion_matrix(y_test,y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap="GnBu")
```


accuracy: 0.9969

```
=====
              precision    recall  f1-score   support

    0.0         1.00      1.00      1.00     12704
    1.0         0.99      1.00      1.00      7296

 accuracy                   1.00      20000
 macro avg              1.00      1.00      1.00      20000
weighted avg              1.00      1.00      1.00      20000
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f8d4050bed0>



Hard Voting Classifier

```
predictions, hard_accuracy, metrics_report, hard_prf = test_model(ensemble_model, X_test,


print('accuracy: {}'.format(hard_accuracy))
print('='*100)
print(metrics_report)
plot_confusion_matrix(ensemble_model, X_test,y_test, xticks_rotation='vertical', cmap="BuP
```

accuracy: 0.986

```
=====
```

	precision	recall	f1-score	support
0.0	1.00	0.98	0.99	12704
1.0	0.97	0.99	0.98	7296
accuracy			0.99	20000
macro avg	0.98	0.99	0.98	20000
weighted avg	0.99	0.99	0.99	20000

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f8d4002b8d0>



Gradient Boosting Classifier:

```
predictions, gb_accuracy, metrics_report, gb_prf = test_model(gb_model, X_test, y_test)

print('accuracy: {}'.format(gb_accuracy))
print('='*100)
print(metrics_report)
plot_confusion_matrix(gb_model, X_test, y_test, xticks_rotation='vertical', cmap="BuPu")
```


accuracy: 0.99855

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	12704
1.0	1.00	1.00	1.00	7296
accuracy			1.00	20000
macro avg	1.00	1.00	1.00	20000
weighted avg	1.00	1.00	1.00	20000

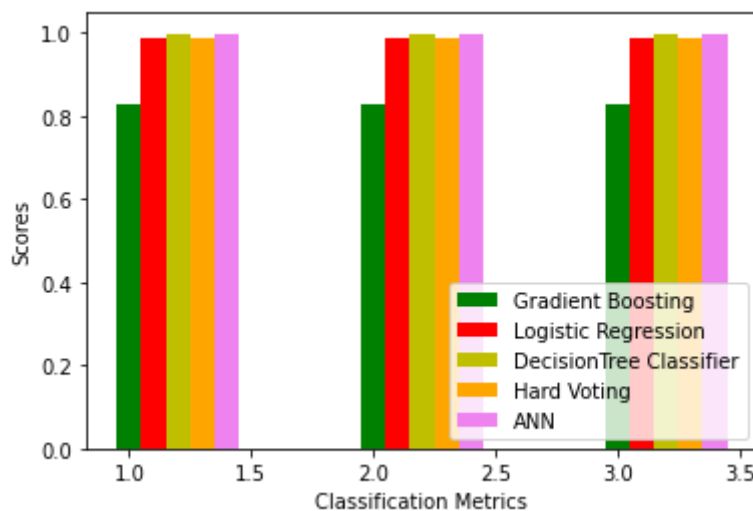
sklearn.metrics.plot_confusion_matrix ConfusionMatrixDisplay at 0x7f58d401e420>

▼ C05 : Model Evaluation

0.0 12677 27

```
# classification matrix indices
xx = np.array([1,2,3])
width = 0.1
gb_prf = np.array([0.83, 0.83, 0.83])
ax = plt.subplot(111)
ax.bar(xx ,height=np.array(gb_prf), width=width, color='g', align='center', label='Gradient
ax.bar(xx + 1*width, height=np.array(log_prf), width=width, color='r', align='center', lab
ax.bar(xx + 2*width, height=np.array(dt_prf), width=width, color='y', align='center', labe
ax.bar(xx + 3*width, height=np.array(hard_prf), width=width, color='orange', align='center
ax.bar(xx + 4*width, height=np.array(ann_prf), width=width, color='violet', align='center'

plt.xlabel('Classification Metrics')
plt.ylabel('Scores')
plt.legend(loc='lower right')
plt.show()
```



```
cl_metric = pd.DataFrame(data = {"Models":[], "Precision":[], "recall":[], "f1-score":[], "acc
prfs=[("Logistic Regression", log_prf), ("Logistic Regression (PCA)", log_pca_prf), ("Support
names=[]
pre=[]
rec=[]
f1=[]
ac=[log_accuracy, log_pca_accuracy, svm_accuracy, dt_accuracy, ann_accuracy, hard_accuracy, gb_a
```

```

for name,(p,r,f) in prfs:
    names.append(name)
    pre.append(p)
    rec.append(r)
    f1.append(f)
cl_metric["Models"]=names
cl_metric["Precision"]=pre
cl_metric["recall"]=rec
cl_metric["f1-score"]=f1
cl_metric["accuracy_score"]=ac

```

```
print(cl_metric.to_string(index=False))
```

Models	Precision	recall	f1-score	accuracy_score
Logistic Regression	0.985110	0.98480	0.984844	0.98480
Logistic Regression (PCA)	0.918738	0.91315	0.910845	0.91315
Support Vector Machine	0.986340	0.98605	0.986089	0.98605
Decision Tree	0.998100	0.99810	0.998100	0.99810
ANN	0.996920	0.99690	0.996902	0.99690
Hard Voting	0.986283	0.98600	0.986039	0.98600
Gradient Boosting	0.830000	0.83000	0.830000	0.99855

Inference: Since we are convinced with the accuracy received from the Boosting model (Gradient) it was not necessary to use the Bagging and Stacking techniques for the prediction

So, from the above predictions it clearly indicates that the Gradient Boosting Model is much reliable in comparison with other models